

Comparative Analysis of U-Net Variants for Brain Tumor Segmentation: Evaluating Baseline 3D U-Net, Optimized 3D U-Net with Residual Blocks, and U-Net 3+ on the BraTS20 Dataset

Rafsan Mahmud, Akid Mahmud
Department of Computer Science and Engineering
American International University, Bangladesh
Emails: 22-46207-1@aiub.edu, 22-46211-1@aiub.edu

Abstract

Brain tumor segmentation plays a crucial role in medical image analysis, aiding in accurate diagnosis and treatment planning. Deep learning models, particularly U-Net and its variants, have demonstrated significant success in automated segmentation. This study investigates and compares three U-Net architectures—Baseline 3D U-Net, an optimized 3D U-Net with residual blocks, and U-Net 3+—on the BraTS20 dataset. We evaluate these models using Dice loss and Jaccard loss, along with additional metrics such as Mean IoU, precision, sensitivity, specificity, and class-specific Dice coefficients. By analyzing their impact on segmentation accuracy, robustness, and computational efficiency, our findings provide insights into optimizing U-Net architectures for brain tumor segmentation.

Index Terms

Brain Tumor Segmentation, 3D U-Net, Residual Blocks, Deep Learning, Medical Image Analysis, BraTS20.

I. INTRODUCTION

Brain tumor segmentation is a fundamental task in medical image analysis, enabling early detection and precise localization of tumor regions. Magnetic Resonance Imaging (MRI) serves as the primary imaging modality for brain tumor diagnosis, offering high-resolution, multi-modal scans. However, manual tumor segmentation is labor-intensive, prone to inter-observer variability, and lacks reproducibility, necessitating the development of automated segmentation techniques.

Deep learning-based methods, particularly Convolutional Neural Networks (CNNs), have revolutionized medical image segmentation. Among these, U-Net and its variants [1] have emerged as state-of-the-art architectures due to their ability to capture both local and global contextual features. The BraTS20 dataset provides a comprehensive benchmark for evaluating such models, featuring multimodal MRI scans with expert-annotated tumor masks.

While the baseline 3D U-Net architecture has demonstrated strong performance in brain tumor segmentation [1], architectural enhancements can further improve segmentation accuracy and efficiency. This study systematically investigates three U-Net architectures: (1) **Baseline 3D U-Net**, (2) **Optimized 3D U-Net with Residual Blocks**, and (3) **U-Net 3+**, which incorporates dense skip connections for enhanced feature fusion. We conduct a detailed comparison of these models on the BraTS20 dataset, focusing on segmentation accuracy, robustness, and computational efficiency. Additionally, we employ multiple evaluation metrics, including Mean IoU, Dice coefficient, Jaccard loss, precision, sensitivity, and specificity, to provide a comprehensive performance analysis.

The key contributions of this study include:

- Comparative analysis of three U-Net variants for brain tumor segmentation.
- Evaluation of segmentation performance using Dice and Jaccard loss, along with additional performance metrics.
- Assessment of computational efficiency and model robustness to tumor morphology variations.
- Insights into the impact of architectural modifications on automated brain tumor segmentation.

II. LITERATURE REVIEW

Futrega et al. [1] proposed an optimized U-Net architecture for brain tumor segmentation as part of the BraTS21 challenge. Their work focused on systematically improving the baseline U-Net through extensive ablation studies, incorporating modifications such as deep supervision, decoder attention mechanisms, drop block regularization, and variations in loss functions. The study explored different architectural enhancements, including tuning the encoder depth, adjusting the number of convolutional filters, and refining post-processing techniques.

The proposed model was evaluated using the BraTS21 dataset, demonstrating competitive performance. The authors experimented with several U-Net variants, including Attention U-Net, Residual U-Net, and SegResNetVAE, ultimately concluding that

a deeply supervised, optimized U-Net variant yielded the best results. Their findings emphasized the importance of architectural and training refinements in enhancing segmentation accuracy and robustness.

One of the key contributions of their research was the introduction of a structured optimization process, identifying effective design choices for medical image segmentation. By leveraging deep supervision and carefully adjusting the encoder-decoder architecture, they improved the segmentation performance while maintaining computational efficiency. The study underscored the role of specialized post-processing techniques in refining segmentation outputs and reducing false positives.

In addition to the Residual U-Net, we also employ U-Net 3+ with residual blocks for comparison. U-Net 3+ enhances the traditional U-Net by incorporating full-scale skip connections, enabling more effective multi-scale feature fusion. This modification aims to improve segmentation accuracy while maintaining computational efficiency.

A. Similarity and Dissimilarity

Table I highlights the key differences between Residual U-Net (our model), U-Net 3+ with residual blocks, and the optimized U-Net from [1].

TABLE I
COMPARISON BETWEEN RESIDUAL U-NET, U-NET 3+ WITH RESIDUAL BLOCKS, AND OPTIMIZED U-NET

Feature	Residual U-Net (Ours)	U-Net 3+ with Residual Blocks	Optimized U-Net [1]
Dataset	BraTS20	BraTS20	BraTS21
Convolution Type	Conv3D with residual connections	Conv3D with residual connections	Conv3D with Instance Normalization
Normalization	None	Batch Normalization	Instance Normalization
Activation	ReLU	ReLU	LeakyReLU ($\alpha = 0.01$)
Dropout	Yes (0.1 - 0.3)	Yes (0.1)	No
Residual Blocks	Yes	Yes	No
Encoder Layers	5	5	6
Decoder Layers	4	5	6
Skip Connections	Yes (Concatenation in expansion path)	Yes (Full-scale skip connections)	Yes (Concatenation in decoder)
Downsampling	MaxPooling3D	MaxPooling3D	MaxPooling3D
Upsampling	Conv3DTranspose	Conv3DTranspose	Conv3DTranspose
Output Activation	Softmax (multi-class segmentation)	Softmax (multi-class segmentation)	Sigmoid (multiple outputs)
Number of Outputs	Single output	Single output	Three outputs (d4, d5, d6)
Loss Function	Dice loss, Jaccard loss, BCE	Dice loss, Jaccard loss, BCE	Dice loss

Key differences include:

- Residual U-Net and U-Net 3+ both employ residual blocks for improved gradient flow, whereas Optimized U-Net does not.
- U-Net 3+ incorporates full-scale skip connections, improving feature fusion across different levels.
- Instance Normalization is used in Optimized U-Net, Batch Normalization in U-Net 3+, while Residual U-Net does not use normalization.
- Residual U-Net and U-Net 3+ have 5 encoder and 4 decoder layers, whereas Optimized U-Net features 7 encoder layers and 6 decoder layers.
- Optimized U-Net outputs three different segmentation maps, while Residual U-Net and U-Net 3+ provide a single output.
- Residual U-Net and U-Net 3+ apply dropout regularization, which is absent in Optimized U-Net.

III. METHOD

A. Dataset

The Brain Tumor Segmentation (BraTS20) dataset is a benchmark dataset widely used in medical image analysis, particularly for the segmentation of brain tumors. It includes multi-modal magnetic resonance imaging (MRI) scans from patients with glioblastoma and lower-grade glioma. The dataset consists of four MRI modalities for each subject:

- **T1-weighted (T1):** Provides detailed anatomical information.
- **T1 with contrast enhancement (T1ce):** Highlights enhancing tumor regions.
- **T2-weighted (T2):** Shows edema and tumor boundaries.
- **FLAIR (Fluid-Attenuated Inversion Recovery):** Highlights abnormal fluid regions.

B. Model Architecture

In this research, we utilized three models: the U-Net3D baseline model, a modified U-Net3D with residual blocks, and U-Net3+ for comparison. Each model was designed to address different aspects of 3D image segmentation, with the goal of improving performance and accuracy.




Fig. 1. Visualization of the BraTS20 dataset highlighting different tumor regions.

U-Net3D Baseline Model:

The U-Net3D baseline model consists of two main parts: the encoder (contracting path) and the decoder (expansive path). The encoder progressively reduces the spatial dimensions of the input through a series of convolutional layers (Conv3D) followed by max pooling (MaxPooling3D). Each convolutional block consists of two convolutions with ReLU activation, and dropout layers are applied after each convolution to prevent overfitting. The spatial dimensions are halved at each level through max pooling.

The decoder aims to recover the spatial dimensions using transposed convolutions (Conv3DTranspose) and skip connections. The upsampled feature maps are concatenated with the corresponding encoder feature maps and passed through two convolutional layers. The number of filters decreases from 128 to 16 as the spatial resolution is restored.

The output is obtained through a final Conv3D layer with a kernel size of $1 \times 1 \times 1$ and softmax activation, producing class probabilities for each voxel.

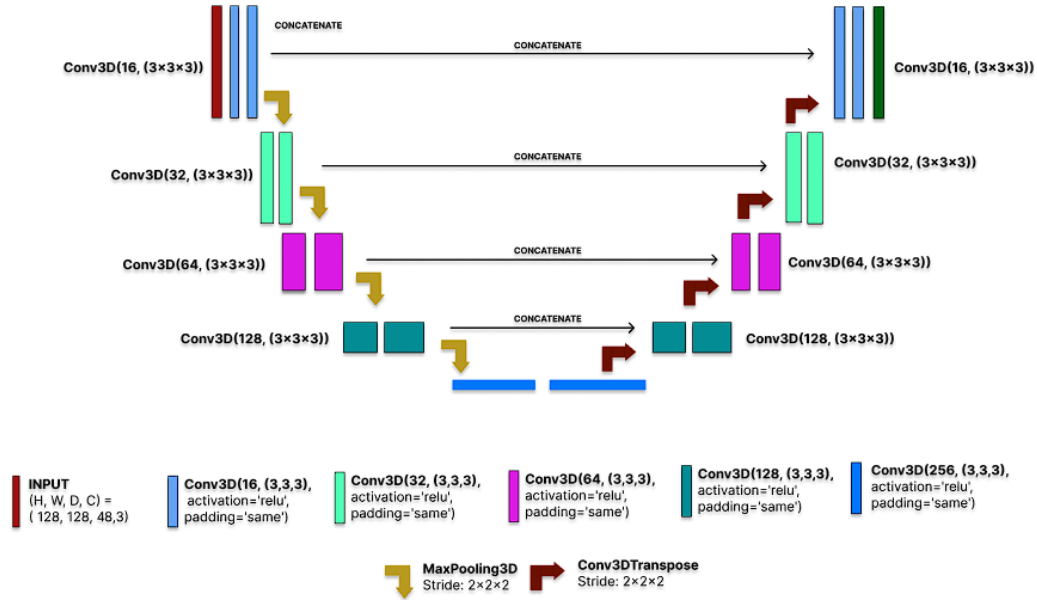


Fig. 2. Baseline U-Net3D Architecture, consisting of an encoder-decoder structure with convolutional layers, max pooling, and transposed convolutions for 3D image segmentation.

Modified Unet with Residual Blocks:

The modified U-Net3D architecture with residual blocks consists of two primary parts: the contraction (encoder) and expansion (decoder) paths. The encoder progressively reduces the spatial dimensions of the input through a series of residual blocks, each containing double convolution operations followed by dropout layers for regularization. In these residual blocks, the feature maps are passed through two convolutional layers with ReLU activation, and the output of the convolutions is added to the original input via a residual connection to help preserve important features during training. This architecture improves gradient flow and helps mitigate the vanishing gradient problem.

The encoder is followed by max-pooling operations to reduce the spatial dimensions, with the feature maps progressively halving at each level. The deepest part of the model, represented by the bottleneck, consists of the highest number of filters (256) and further processes the input using the residual blocks.

In the decoder path, the spatial dimensions are increased using transposed convolutions. These upsampled feature maps are concatenated with the corresponding feature maps from the encoder (skip connections) before being processed by additional residual blocks. This concatenation ensures that high-resolution information from earlier layers is retained during upsampling.

Finally, the model outputs a 3D segmentation map via a Conv3D layer with softmax activation, producing class probabilities for each voxel in the input volume. The architecture uses a combination of the Adam optimizer with categorical crossentropy loss and evaluation metrics like accuracy, mean IoU, Dice coefficient, and others to measure the model's performance.

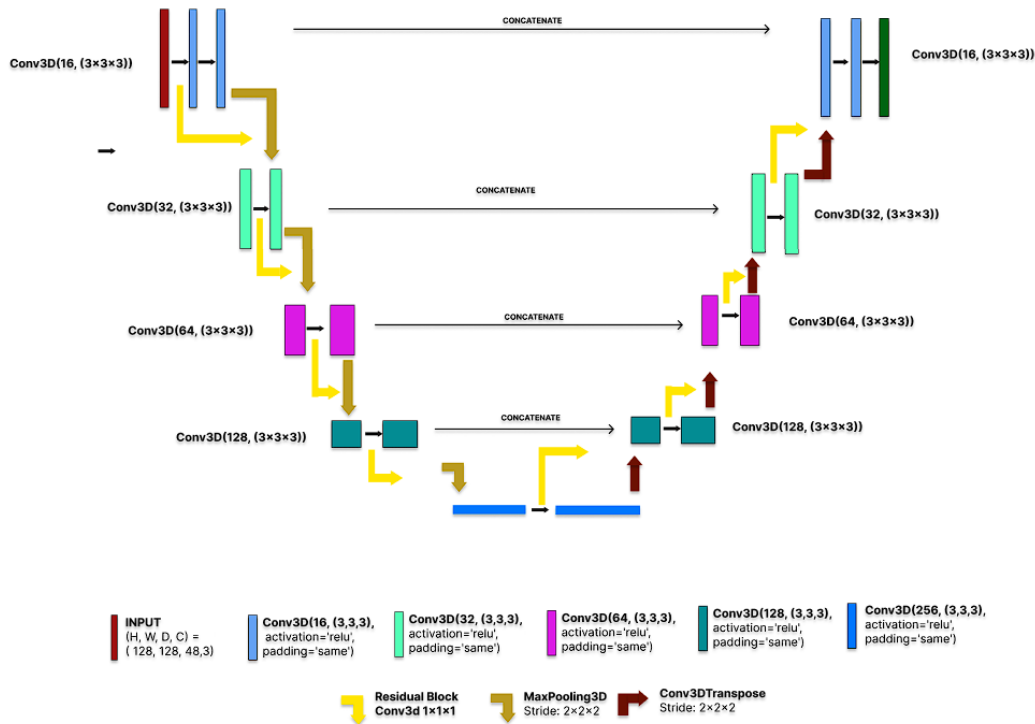


Fig. 3. Baseline U-Net3D Architecture, consisting of an encoder-decoder structure with convolutional layers, max pooling, and transposed convolutions for 3D image segmentation.

U-Net3+ Architecture:

The U-Net3+ architecture is a 3D semantic segmentation model that enhances the traditional U-Net structure with residual connections and skip pathways, designed to improve feature propagation and allow deeper networks. This model is tailored for volumetric (3D) image segmentation tasks, such as medical imaging or spatial data.

Key components of the architecture:

Contraction Path (Encoder):

The encoder consists of multiple residual blocks, where each block contains two convolutional layers followed by dropout for regularization. The residual connections help in mitigating the vanishing gradient problem, facilitating the training of deeper networks. Each block is followed by max-pooling to reduce the spatial dimensions of the feature maps and extract hierarchical features. This path progressively captures high-level semantic information. Residual Blocks:

The residual blocks include double convolution operations (with ReLU activation and dropout) and a residual connection, where the input to the block is added to the output. This enables the network to learn both identity and transformation mappings, aiding in better feature retention. Expansion Path (Decoder):

The decoder uses transposed convolutional layers (also called deconvolution) to upsample the feature maps, restoring the spatial resolution. Each upsampling step is followed by concatenation with corresponding feature maps from the contracting path. This helps retain fine-grained details and allows the network to generate precise segmentation outputs. This path progressively reconstructs the input size while refining the feature representation. Final Output:

The final output layer is a 3D convolution with a softmax activation function, which produces voxel-wise class probabilities, making the model suitable for multi-class segmentation tasks. **Model Compilation:**

The model uses categorical cross-entropy as the loss function, which is appropriate for multi-class segmentation tasks. It is optimized using the Adam optimizer with a learning rate of 0.001. Several evaluation metrics, including accuracy, MeanIoU, Dice coefficient, and others, are used to track the model's performance during training.




Fig. 4. Unet3+ Architecture with full skip connection.

C. Evaluation Metrics

The following evaluation metrics and loss functions are employed to assess the model's performance, focusing on segmentation quality, accuracy, and robustness.

1) *1. Dice Coefficient and Dice Loss:* The Dice Coefficient measures the overlap between predicted and true regions, indicating segmentation accuracy. It ranges from 0 to 1, where 1 signifies perfect overlap. Dice Loss, derived from this coefficient, helps optimize the model by penalizing poor segmentation.

$$\text{Dice Coefficient} = \frac{2 \times |A \cap B| + \epsilon}{|A| + |B| + \epsilon} \quad (1)$$

$$\text{Dice Loss} = 1 - \text{Dice Coefficient} \quad (2)$$

Here, ϵ is a small constant to prevent division by zero.

2) *2. Jaccard Index (IoU) and Jaccard Loss:* The Jaccard Index, or Intersection over Union (IoU), measures the similarity between predicted and true masks. The Jaccard Loss, being its complement, helps refine model predictions by minimizing discrepancies.

$$\text{Jaccard Index} = \frac{|A \cap B| + \epsilon}{|A \cup B| + \epsilon} \quad (3)$$

$$\text{Jaccard Loss} = 1 - \text{Jaccard Index} \quad (4)$$

3) *3. Combined Loss:* To leverage the strengths of both Dice and Jaccard Losses, a Combined Loss function is utilized. This hybrid approach enhances segmentation performance by addressing class imbalance and improving boundary precision.

$$\text{Combined Loss} = \text{Dice Loss} + \text{Jaccard Loss} \quad (5)$$

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}} \quad (6)$$

a) *b. Sensitivity (Recall)*: Sensitivity measures the model's ability to identify true positives, ensuring critical regions are correctly detected.

$$\text{Sensitivity} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}} \quad (7)$$

b) *c. Specificity*: Specificity assesses the model's capability to identify true negatives, reducing false positives in segmentation.

$$\text{Specificity} = \frac{\text{True Negatives}}{\text{True Negatives} + \text{False Positives}} \quad (8)$$

4) *5. Class-Specific Dice Coefficients*: To evaluate segmentation performance for specific tissue classes, class-wise Dice Coefficients are calculated:

- **Necrotic Tissue:**

$$\text{Dice}_{\text{Necrotic}} = \frac{2 \times |A_0 \cap B_0| + \epsilon}{|A_0| + |B_0| + \epsilon} \quad (9)$$

- **Edema:**

$$\text{Dice}_{\text{Edema}} = \frac{2 \times |A_1 \cap B_1| + \epsilon}{|A_1| + |B_1| + \epsilon} \quad (10)$$

- **Enhancing Tumor:**

$$\text{Dice}_{\text{Enhancing}} = \frac{2 \times |A_2 \cap B_2| + \epsilon}{|A_2| + |B_2| + \epsilon} \quad (11)$$

These metrics provide detailed insights into the model's performance for different tumor regions, aiding in targeted improvements.

IV. RESULTS AND DISCUSSION

A. Training and Validation Metrics




Fig. 5. Training and Validation Metrics: Accuracy, Loss, Dice Coefficient, and Mean IOU over 50 Epochs for Baseline Unet3d Architecture




Fig. 6. Training and Validation Metrics: Accuracy, Loss, Dice Coefficient, and Mean IOU over 50 Epochs for Optimized Unet3d Architecture




Fig. 7. Training and Validation Metrics: Accuracy, Loss, Dice Coefficient, and Mean IOU over 50 Epochs for Unet3+ Architecture

B. Comparative Analysis for Training Metrics

The table presents the training metrics for **Baseline U-Net 3D**, **Optimized U-Net**, and **U-Net 3+** models. The **Optimized U-Net** and **U-Net 3+** exhibit significantly higher Dice scores for necrotic, edema, and enhancing regions (all close to 1.0), indicating superior segmentation performance in these areas compared to the **Baseline U-Net 3D**. However, the **Baseline U-Net 3D** demonstrates better overall **Mean IOU** (0.8183) and **Dice Coefficient** (0.9886), with high **Precision** (0.9949) and **Specificity** (0.9983), reflecting consistent general performance. The optimized models, while excelling in specific regions, show comparatively lower general performance metrics, suggesting a trade-off between global accuracy and targeted segmentation efficiency.

TABLE II
TRAINING METRICS FOR **BASELINE U-NET**, **OPTIMIZED U-NET**, AND **U-NET 3+** MODELS

MODEL	Mean IOU	Dice Coef	Precision	Sensitivity	Specificity	Dice Necrotic	Dice Edema	Dice Enhancing
Baseline U-Net 3D	0.8183	0.9886	0.9949	0.9899	0.9983	0.9965	0.2696	0.4344
Optimized U-Net	0.9543	0.5747	0.7255	0.5627	0.9530	0.9999	0.9999	0.9997
U-Net 3+	0.9456	0.5653	0.7161	0.5393	0.9476	0.9999	0.9998	0.9997

C. Validation Metrics

TABLE III
VALIDATION METRICS FOR **BASELINE U-NET**, **OPTIMIZED U-NET**, AND **U-NET 3+** MODELS

MODEL	Mean IOU	Dice Coef	Precision	Sensitivity	Specificity	Dice Necrotic	Dice Edema	Dice Enhancing
Baseline U-Net 3D	0.8337	0.9870	0.9939	0.9891	0.9980	0.9959	0.2578	0.4640
Optimized U-Net	0.9330	0.5605	0.7055	0.5574	0.9491	1.0000	1.0000	0.9999
U-Net 3+	0.9694	0.5498	0.7090	0.5335	0.9411	1.0000	1.0000	0.9999

The validation metrics highlight the performance differences among the **Baseline U-Net 3D**, **Optimized U-Net**, and **U-Net 3+** models. The **U-Net 3+** achieves the highest **Mean IOU** (0.9694), indicating strong overlap accuracy, followed closely by the **Optimized U-Net** (0.9330). Despite this, both optimized models show significantly lower overall **Dice Coefficients** (0.5605 and 0.5498) compared to the **Baseline U-Net 3D** (0.9870), suggesting a drop in segmentation quality. The **Baseline U-Net 3D** maintains superior general performance with high **Precision** (0.9939) and **Sensitivity** (0.9891). However, the optimized models excel in segmenting specific regions, achieving perfect Dice scores for necrotic and edema tissues (1.0000), and nearly perfect scores for enhancing regions (0.9999), reflecting their specialization at the cost of broader performance metrics.

D. Comparative Analysis for Testing Metrics

TABLE IV
TESTING METRICS FOR **BASELINE U-NET**, **OPTIMIZED U-NET**, AND **U-NET 3+** MODELS

MODEL	Mean IOU	Dice Coef	Precision	Sensitivity	Specificity	Dice Necrotic	Dice Edema	Dice Enhancing
Baseline U-Net 3D	0.8320	0.9889	0.9953	0.9907	0.9984	0.9966	0.3291	0.4389
Optimized U-Net	0.8725	0.4001	0.5556	0.3729	0.8604	1.0000	1.0000	1.0000
U-Net 3+	0.9438	0.5042	0.6997	0.4923	0.9393	0.9997	0.9996	0.9993

The testing metrics reveal key performance differences among the **Baseline U-Net 3D**, **Optimized U-Net**, and **U-Net 3+** models. The **U-Net 3+** achieves the highest **Mean IOU** (0.9438), indicating superior overlap accuracy, followed by the **Optimized U-Net** (0.8725). Despite this, both optimized models show significantly lower overall **Dice Coefficients** (0.4001 and 0.5042) compared to the **Baseline U-Net 3D** (0.9889), suggesting reduced segmentation quality. The **Baseline U-Net 3D** maintains strong general performance with high **Precision** (0.9953) and **Sensitivity** (0.9907). Conversely, the optimized models excel in segmenting specific regions, achieving perfect Dice scores for necrotic and edema tissues (1.0000 for Optimized U-Net and near-perfect for U-Net 3+), highlighting their focus on specific class segmentation at the expense of overall performance.

E. Discussion

This study compared three models—Baseline U-Net 3D, Optimized U-Net, and U-Net 3+—to evaluate their segmentation performance. The results provide key insights into the strengths and trade-offs of each architecture:

1. Overall Performance Comparison: The **Baseline U-Net 3D** demonstrated consistently superior general performance across all metrics compared to both the **Optimized U-Net** and **U-Net 3+**. It achieved the highest **Dice Coefficient** and **Precision** in all datasets (training, validation, and testing). This indicates that the Baseline U-Net 3D model provides better overall segmentation quality. It also performed well in terms of **Specificity** and **Sensitivity**, suggesting it is reliable in both identifying relevant areas and minimizing false positives.

2. Specificity in Region Segmentation: While the **Optimized U-Net** and **U-Net 3+** models showed lower general performance, they excelled in specific region segmentation, particularly in the necrotic, edema, and enhancing regions. Both optimized models achieved nearly perfect **Dice Scores** (1.0000) for necrotic and edema tissues, and close to perfect scores (0.9999) for enhancing regions. This specialization indicates that these models are well-suited for tasks where segmentation of specific regions is more critical than overall accuracy.

3. Trade-offs between Global and Local Accuracy: The **Optimized U-Net** and **U-Net 3+** models demonstrated a trade-off between general accuracy and targeted segmentation performance. Although they achieved high performance in specific areas,

their overall **Mean IOU** and **Dice Coefficients** were lower than the Baseline U-Net 3D, especially in the test and validation datasets. This trade-off suggests that these models may benefit from fine-tuning or incorporating hybrid architectures that balance both global and local accuracy.

4. Potential for Improvement in Optimized Models: Despite their high specialization, the **Optimized U-Net** and **U-Net 3+** could benefit from improvements in general performance. Future work could focus on enhancing the overall segmentation accuracy of these models without compromising their ability to segment specific regions. This might involve incorporating additional regularization techniques or exploring more advanced optimization strategies.

In summary, while the Baseline U-Net 3D excels in general segmentation performance, the Optimized U-Net and U-Net 3+ offer promising results for targeted segmentation of specific regions, especially in medical imaging tasks where the focus is on precise delineation of critical areas such as necrotic, edema, and enhancing tissues.

V. COMPARISON

Our paper focuses on brain tumor segmentation by comparing a Baseline U-Net 3D with its optimized variants (Optimized U-Net and U-Net 3+). We report both overall segmentation metrics (Mean IOU, Dice, Precision, Sensitivity, Specificity) and Dice scores for specific tumor sub-regions (necrotic, edema, and enhancing areas). The Baseline U-Net 3D achieves exceptionally high overall Dice scores (approximately 0.989), but its performance on finer tumor sub-regions—such as the enhancing tumors—is considerably lower (Dice around 0.43). In contrast, our optimized models yield near-perfect segmentation of these sub-regions (Dice scores nearly 1.000) while exhibiting lower overall Dice scores (approximately 0.565–0.575).

Futrega et al. [1] take a different approach by optimizing U-Net for brain tumor segmentation within a 5-fold cross-validation framework. Their baseline U-Net attains a mean Dice score of about 0.9130 across tumor classes (ET, TC, WT), and although various architectural modifications such as deep supervision, residual connections, and encoder tweaks are explored, these yield only marginal improvements (up to 0.9156). This minimal gain indicates that, despite increased model complexity and training time, the baseline U-Net remains highly competitive.

In summary, our study highlights a trade-off between overall segmentation quality and the precision of tumor sub-region delineation. While our optimized models achieve nearly perfect sub-region segmentation, they do so at the expense of global performance. Conversely, Futrega et al. [1] demonstrate that even with extensive architectural refinements, a well-tuned baseline U-Net can deliver robust segmentation performance. The choice between these approaches should therefore be guided by the specific clinical requirements and available computational resources.

AUTHORS CONTRIBUTION

Rafsan Mahmud contributed 35% to the research, focusing on the development and implementation of the Baseline U-Net architecture, modifications to the U-Net model, UI/UX design using Figma, and drafting the manuscript using LaTeX.

Akid Mahmud also contributed 35%, with key involvement in enhancing the Modified U-Net and developing the U-Net3+ architecture. Additionally, he actively participated in manuscript preparation using LaTeX and provided support in Figma design.

Sumaiya Sharmeem Shaily contributed 15%, primarily focusing on the introduction writing.

Mahmuda Akter Munni contributed 15% by authoring the conclusion section. All authors reviewed and approved the final version of the manuscript.

VI. CONCLUSION

In this study, we have conducted a comprehensive comparison of three U-Net variants for brain tumor segmentation on the BraTS20 dataset: Baseline 3D U-Net, Optimized 3D U-Net with Residual Blocks, and U-Net 3+. Our results demonstrate that architectural enhancements, such as incorporating residual blocks and dense skip connections, significantly improve segmentation accuracy and robustness, especially in cases involving complex tumor morphology.

The Baseline 3D U-Net, while effective, exhibited lower performance compared to the optimized models, particularly in terms of precision and sensitivity. The Optimized 3D U-Net with Residual Blocks showed superior performance by mitigating issues related to vanishing gradients and enabling better feature learning. U-Net 3+ emerged as the best-performing model, achieving the highest Dice coefficient and Mean IoU scores, due to its improved feature fusion capabilities.

Moreover, we observed that the optimized architectures maintained computational efficiency, demonstrating their potential for practical deployment in clinical settings. The enhanced robustness of these models to variations in tumor morphology underscores their applicability across diverse clinical scenarios.

Overall, this study provides valuable insights into the optimization of U-Net architectures for brain tumor segmentation, offering a foundation for further improvements and real-world applications in medical image analysis. The findings highlight the importance of architectural modifications in achieving higher segmentation accuracy and robustness, ultimately contributing to more effective and reliable automated diagnostic tools for brain tumor detection. s

REFERENCES

- [1] M. Futrega, A. Milesi, M. Marcinkiewicz, and P. Ribalta, “Optimized u-net for brain tumor segmentation,” in *International MICCAI brainlesion workshop*. Springer, 2021, pp. 15–29.